Econ 103 – Statistics for Economists

Chapter 2 feat. Z Scores and OLS

Mallick Hossain

University of Pennsylvania

For next time!
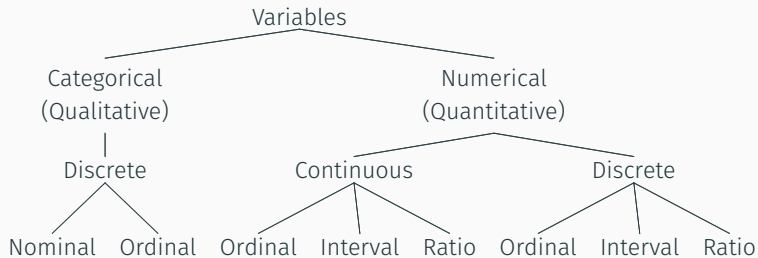
Types of Variables

- What are the differences between the following variables?
  - "Age" and "gender"
  - "Gender" and "class standing"
  - "SAT score" and "job creation"

Variables

Categorical (Qualitative) — Numerical (Quantitative)

Discrete

Nominal, Ordinal

Continuous — Ordinal, Interval, Ratio

Discrete — Ordinal, Interval, Ratio

- **Discrete:** Can be a countable number of values
- **Continuous:** Can take on any value

Can you order the following from weakest to strongest?
Interval, Nominal, Ordinal, Ratio.

- **Nominal:** no order to the categories

- **Nominal:** no order to the categories
- **Ordinal:** categories with natural order

- **Nominal:** no order to the categories
- **Ordinal:** categories with natural order
- **Interval:** only differences meaningful, no natural zero

- **Nominal:** no order to the categories
- **Ordinal:** categories with natural order
- **Interval:** only differences meaningful, no natural zero
- **Ratio:** differences and ratios meaningful, natural zero

Summary Statistics

1. Measures of Central Tendency
   - **Mean:** the average ("balance point")

$$\bar{x} = \frac{1}{n} \sum_{i=1}^{n} x_i$$

   - **Median:** the middle observation (if data has even number of observations, take the mean of the middle two observations)

2. Percentiles

P^{th} Percentile = Value in $(P/100) \cdot (n+1)^{th}$ Ordered Position

60 63 65 67 70 72 75 75 80 82 84 85

$$
\begin{aligned}
Q_1 &= \text{value in the } 0.25(n+1)^{th} \text{ ordered position} \\
&= \text{value in the } 3.25^{th} \text{ ordered position} \\
&= 0.75 * 65 + 0.25 * 67 \\
&= 65.5
\end{aligned}
$$

Definitions

3. Measures of Spread
   - **Variance:** the spread from the mean

   $$s^2 = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2$$

   - **Standard Deviation:** another way to measure the spread

   $$s = \sqrt{s^2}$$

   - **Range:** the distance between the highest and lowest value

   $$Range = |x_{max} - x_{min}|$$

   - **Interquartile Range (IQR):** the distance between the upper and lower quartiles

   $$IQR = |x_{75\%} - x_{25\%}|$$

Why Squares?

$$s^2 = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2$$

$$s^2 = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2$$

What's Wrong With This?

$$\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x}) = \frac{1}{n-1} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} \bar{x} \right] = \frac{1}{n-1} \left[\sum_{i=1}^{n} x_i - n\bar{x} \right]$$

$$= \frac{1}{n-1} \left[\sum_{i=1}^{n} x_i - n \cdot \frac{1}{n} \sum_{i=1}^{n} x_i \right]$$

$$= \frac{1}{n-1} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} x_i \right] = 0$$

4. Measure of Symmetry
   - **Skewness:** a measure of symmetry, positive values means the right tail is longer and vice versa

$$Skewness = \frac{1}{n} \frac{\sum_{i=1}^{n}(x_i - \bar{x})^3}{s^3}$$

Skewness – A Measure of Symmetry

$$\text{Skewness} = \frac{1}{n} \frac{\sum_{i=1}^{n}(x_i - \bar{x})^3}{s^3}$$

What do the values indicate?

Zero $\Rightarrow$ symmetry, positive right-skewed, negative left-skewed.

Why cubed?

To get the desired sign.

Why divide by s^3?

So that skewness is unitless

Rule of Thumb

Typically (but not always), right-skewed $\Rightarrow$ mean $>$ median
left-skewed $\Rightarrow$ mean $<$ median

5. Relationship between variables (to be covered later)
   - **Covariance:** how two variables vary together
   - **Correlation:** normalized version of covariance (can only range between -1 and +1)
   - **Regression:** an estimation of how two variables are related
   - We'll cover these in-depth later in the semester

Charts

Some Data Visualization Quotes

1. "Overload, clutter, and confusion are not attributes of information, they are failures of design" –Edward Tufte

2. "...few people will appreciate the music if I just show them the notes. Most of us need to listen to the music to understand how beautiful it is. But often, that's how we present statistics; we just show the notes we don't play the music." –Hans Rosling

```
https://i.imgur.com/W4BKCVU.gif
```

Before

Calories per 100g for different foods

Legend:
- French Fries
- Potato Chips
- Bacon
- Pizza
- Chili Dog

Number of Calories (y-axis): 0, 100, 200, 300, 400, 500, 600, 700

Type of Food (x-axis): French Fries, Potato Chips, Bacon, Pizza, Chili Dog

Created by **Darkhorse Analytics**   www.**darkhorseanalytics**.com

After

Calories per 100g

607	542	533	296	260
French Fries	Potato Chips	Bacon	Pizza	Chili Dog

Created by **Darkhorse Analytics**   www.**darkhorseanalytics**.com

Illustrating with Charts (Box and Whisker Chart)

Salaries by Department in Chicago

Illustrating with Charts (Histogram)

Salaries in the Chicago Police Department

1. Histograms show the frequency of different observations
2. Important Choice: How many bins?

Too Few Bins (Oversmoothing)

Salaries in the Chicago Police Department

Too Many Bins (Undersmoothing)

Salaries in the Chicago Police Department

Just Right! (Usually around 20 bins or so)

Salaries in the Chicago Police Department

1. What does it measure?
2. What are its units compared to those of the data?
3. How do its units change if those of the data change?
4. What are the benefits and drawbacks of this statistic?

Some of the information regarding items 2 and 3 is on the homework rather than in the slides because working it out for yourself is a good way to check your understanding.

Outliers

Outlier: A very unusual observation relative to the other observations in the dataset (i.e. very small or very big).

Which Summary Stats are Sensitive to Outliers?

- Assume our data is 1, 2, 3, 4, 5. What are our summary stats (mean, median, variance, range, IQR)
- What will be affected if the data includes an outlier and becomes 1, 2, 3, 4, 4990?

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3
- Variance changes from 2.5 to 4,975,032

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3
- Variance changes from 2.5 to 4,975,032
- Range changes from 4 to 4889

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3
- Variance changes from 2.5 to 4,975,032
- Range changes from 4 to 4889
- IQR remains at 2

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3
- Variance changes from 2.5 to 4,975,032
- Range changes from 4 to 4889
- IQR remains at 2
- When Does the Median Change? IQR?

Which Summary Stats are Sensitive to Outliers?

- Mean changes from 3 to 1000
- Median remains at 3
- Variance changes from 2.5 to 4,975,032
- Range changes from 4 to 4889
- IQR remains at 2
- When Does the Median Change? IQR?
    - Ranks would have to change.

Summary of Sensitivity

Variance
Essentially the average squared distance from the mean.
Sensitive to both skewness and outliers.

Standard Deviation
$\sqrt{\text{Variance}}$, but more convenient since same units as data

Range
Difference between larges and smallest observations. *Very*
sensitive to outliers.

Interquartile Range
Range of middle 50% of the data. Insensitive to outliers,
skewness.

Sample vs. Population

For now, you can think of the population as a list of N objects:

Population: $x_1, x_2, \ldots, x_N$

from which we draw a sample of size $n < N$ objects:

Sample: $x_1, x_2, \ldots, x_n$

Important Point:

Later in the course we'll be more formal by considering probability models that represent the *act of sampling* from a population rather than thinking of a population as a list of objects. Once we do this we will no longer use the notation N as the population will be *conceptually infinite*.

Essential Distinction: Parameter vs. Statistic

N individuals in the Population, *n* individuals in the Sample:

	Parameter (Population)	Statistic (Sample)
Mean	$\mu = \dfrac{1}{N} \sum\limits_{i=1}^{N} x_i$	$\bar{x} = \dfrac{1}{n} \sum\limits_{i=1}^{n} x_i$
Var.	$\sigma^2 = \dfrac{1}{N} \sum\limits_{i=1}^{N} (x_i - \mu)^2$	$s^2 = \dfrac{1}{n-1} \sum\limits_{i=1}^{n} (x_i - \bar{x})^2$
S.D.	$\sigma = \sqrt{\sigma^2}$	$s = \sqrt{s^2}$

Key Point

We use a sample $x_1, \ldots, x_n$ to calculate statistics (e.g. $\bar{x}$, s^2, s) that serve as estimates of the corresponding population parameters (e.g. μ, σ^2, σ).

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^{N} (x_i - \mu)^2$$

$$s^2 = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2$$

There is an important reason for this, but explaining it requires some concepts we haven't learned yet.

Some Intuition

- **Intuition 1:** If we only had one data point, what would be the sample variance? Would it even be defined?
- **Intuition 2:** We know that the deviations from the sample mean sum to zero (see discussion of why variance is squared). Hence, we only need to know $n-1$ of the deviations since the last one will be whatever it takes to make the sum of them equal to 0. Hence, it would be proper to divide by $n-1$ instead of n

Z Scores

Empirical Rule

For large populations that are approximately bell-shaped, standard deviation tells where most observations will be relative to the mean:

- $\approx$ 68% of observations are in the interval $\mu \pm \sigma$
- $\approx$ 95% of observations are in the interval $\mu \pm 2\sigma$
- Almost all of observations are in the interval $\mu \pm 3\sigma$

Z-scores: How many standard deviations from the mean?

$$z_i = \frac{x_i - \bar{x}}{s}$$

$$z_i = \frac{x_i - \bar{x}}{s}$$

Unitless

Allows comparison of variables with different units.

Z-scores: How many standard deviations from the mean?

$$z_i = \frac{x_i - \bar{x}}{s}$$

Unitless
Allows comparison of variables with different units.

Detecting Outliers
Measures how "extreme" one observation is relative to the others.

Z-scores: How many standard deviations from the mean?

$$z_i = \frac{x_i - \bar{x}}{s}$$

Unitless
Allows comparison of variables with different units.

Detecting Outliers
Measures how "extreme" one observation is relative to the others.

Linear Transformation

What is the sample mean of the z-scores?

What is the sample mean of the z-scores?

$$\bar{z} = \frac{1}{n} \sum_{i=1}^{n} z_i = \frac{1}{n} \sum_{i=1}^{n} \frac{x_i - \bar{x}}{s} = \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} \bar{x} \right]$$

What is the sample mean of the z-scores?

$$\bar{z} = \frac{1}{n} \sum_{i=1}^{n} z_i = \frac{1}{n} \sum_{i=1}^{n} \frac{x_i - \bar{x}}{s} = \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} \bar{x} \right]$$

$$= \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - n\bar{x} \right] = \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - n \cdot \frac{1}{n} \sum_{i=1}^{n} x_i \right]$$

What is the sample mean of the z-scores?

$$\bar{z} = \frac{1}{n} \sum_{i=1}^{n} z_i = \frac{1}{n} \sum_{i=1}^{n} \frac{x_i - \bar{x}}{s} = \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} \bar{x} \right]$$

$$= \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - n\bar{x} \right] = \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - n \cdot \frac{1}{n} \sum_{i=1}^{n} x_i \right]$$

$$= \frac{1}{n \cdot s} \left[\sum_{i=1}^{n} x_i - \sum_{i=1}^{n} x_i \right] = 0$$

What is the variance of the z-scores?

What is the variance of the z-scores?

$$s_z^2 = \frac{1}{n-1} \sum_{i=1}^{n} (z_i - \bar{z})^2 = \frac{1}{n-1} \sum_{i=1}^{n} z_i^2 = \frac{1}{n-1} \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s_x} \right)^2$$

What is the variance of the z-scores?

$$s_z^2 = \frac{1}{n-1} \sum_{i=1}^{n} (z_i - \bar{z})^2 = \frac{1}{n-1} \sum_{i=1}^{n} z_i^2 = \frac{1}{n-1} \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s_x} \right)^2$$

$$= \frac{1}{s_x^2} \left[\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2 \right] = \frac{s_x^2}{s_x^2} = 1$$

$$s_z^2 = \frac{1}{n-1} \sum_{i=1}^{n} (z_i - \bar{z})^2 = \frac{1}{n-1} \sum_{i=1}^{n} z_i^2 = \frac{1}{n-1} \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s_x} \right)^2$$

$$= \frac{1}{s_x^2} \left[\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2 \right] = \frac{s_x^2}{s_x^2} = 1$$

So what is the *standard deviation* of the z-scores?

If we knew the population mean μ and standard deviation σ we could create a *population version* of a z-score. This leads to an important way of rewriting the Empirical Rule: Bell-shaped

population $\Rightarrow$ approx. 95% of observations x_i satisfy

$$\mu - 2\sigma \leq x_i \leq \mu + 2\sigma$$

$$-2\sigma \leq x_i - \mu \leq 2\sigma$$

$$-2 \leq \frac{x_i - \mu}{\sigma} \leq 2$$

Covariance and Correlation

Two Samples of Numeric Data

$x_1, \ldots, x_n$ and $y_1, \ldots, y_n$

Dependence

Do *x* and *y* both tend to be large (or small) at the same time?

Key Point

Use the idea of centering and standardizing to decide what "big" or "small" means in this context.

Are Engine Cylinders and Horsepower Related?

Horsepower and MPG

Recall Formulas

$$\bar{x} = \frac{1}{n} \sum_{i=1}^{n} x_i$$

$$\bar{y} = \frac{1}{n} \sum_{i=1}^{n} y_i$$

$$s_x = \sqrt{\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2}$$

$$s_y = \sqrt{\frac{1}{n-1} \sum_{i=1}^{n} (y_i - \bar{y})^2}$$

Covariance

$$s_{xy} = \frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})(y_i - \bar{y})$$

- Centers each observation around its mean and multiplies.
- Zero $\Rightarrow$ no linear dependence
- Positive $\Rightarrow$ positive linear dependence
- Negative $\Rightarrow$ negative linear dependence
- Population parameter: σ_{xy}
- Units?

Correlation

$$r_{xy} = \frac{1}{n-1} \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s_x} \right) \left(\frac{y_i - \bar{y}}{s_y} \right) = \frac{s_{xy}}{s_x s_y}$$

- Centers *and* standardizes each observation
- Bounded between -1 and 1
- Zero $\Rightarrow$ no linear dependence
- Positive $\Rightarrow$ positive linear dependence
- Negative $\Rightarrow$ negative linear dependence
- Population parameter: ρ_{xy}
- Unitless

guessthecorrelation.com

Introduction to Regression

Horsepower and MPG

- In order to fit a line through this, we need to estimate
  $y = a + bx$
- How do we find a and b?

- Linear regression chooses the slope (*b*) and intercept (*a*) that minimize the sum of squared vertical deviations

$$\sum_{i=1}^{n} d_i^2 = \sum_{i=1}^{n}(y_i - a - bx_i)^2$$

- Why do we square the deviations?

$$\underset{a,b}{\text{minimize}} \sum_{i=1}^{n} d_i^2 = \sum_{i=1}^{n}(y_i - a - bx_i)^2$$

$$\hat{y} = a + bx$$

- $(x_i, y_i)_{i=1}^{n}$ are the observed data
- $\hat{y}$ is our prediction for a given value of x
- Neither x nor $\hat{y}$ needs to be in out dataset!

$$\widehat{hp} = 324.08 - 8.83mpg$$

$$\widehat{hp} = 324.08 - 8.83mpg$$

$$76.84 = 324.08 - 8.83 * 28$$

Minimize the sum of squared vertical deviations from the line:

$$\min_{a,b} \sum_{i=1}^{n} (y_i - a - bx_i)^2$$

How should we proceed? Select all that apply.

(a) Differentiate with respect to x

(b) Differentiate with respect to y

(c) Differentiate with respect to a

(d) Differentiate with respect to b

(e) You can't fool me! You can't solve this with calculus.

Objective Function

$$\min_{a,b} \sum_{i=1}^{n} (y_i - a - bx_i)^2$$

FOC with respect to a

Derivation

Objective Function

$$\min_{a,b} \sum_{i=1}^{n} (y_i - a - bx_i)^2$$

FOC with respect to a

$$-2 \sum_{i=1}^{n} (y_i - a - bx_i) = 0$$

$$\sum_{i=1}^{n} y_i - \sum_{i=1}^{n} a - b \sum_{i=1}^{n} x_i = 0$$

$$\frac{1}{n} \sum_{i=1}^{n} y_i - \frac{na}{n} - \frac{b}{n} \sum_{i=1}^{n} x_i = 0$$

$$\bar{y} - a - b\bar{x} = 0$$

$$\bar{y} = a + b\bar{x}$$

Substitute $a = \bar{y} - b\bar{x}$

$$\sum_{i=1}^{n}(y_i - a - bx_i)^2 = \sum_{i=1}^{n}(y_i - \bar{y} + b\bar{x} - bx_i)^2$$

$$= \sum_{i=1}^{n}\left[(y_i - \bar{y}) - b(x_i - \bar{x})\right]^2$$

FOC wrt b

$$-2\sum_{i=1}^{n}\left[(y_i - \bar{y}) - b(x_i - \bar{x})\right](x_i - \bar{x}) = 0$$

$$\sum_{i=1}^{n}(y_i - \bar{y})(x_i - \bar{x}) - b\sum_{i=1}^{n}(x_i - \bar{x})^2 = 0$$

$$b = \frac{\sum_{i=1}^{n}(y_i - \bar{y})(x_i - \bar{x})}{\sum_{i=1}^{n}(x_i - \bar{x})^2}$$

Simple Linear Regression

Problem

$$\min_{a,b} \sum_{i=1}^{n} (y_i - a - bx_i)^2$$

Solution

$$b = \frac{\sum_{i=1}^{n} (y_i - \bar{y})(x_i - \bar{x})}{\sum_{i=1}^{n} (x_i - \bar{x})^2}$$

$$a = \bar{y} - b\bar{x}$$

Relating Regression to Covariance and Correlation

$$b = \frac{\sum_{i=1}^{n} (y_i - \bar{y})(x_i - \bar{x})}{\sum_{i=1}^{n} (x_i - \bar{x})^2} = \frac{\frac{1}{n-1} \sum_{i=1}^{n} (y_i - \bar{y})(x_i - \bar{x})}{\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \bar{x})^2} = \frac{s_{xy}}{s_x^2}$$

$$r = \frac{s_{xy}}{s_x s_y} = b \frac{s_x}{s_y}$$

Units

Correlation is unitless, covariance and regression coefficients (a, b) are not. (What are the units of these?)

Symmetry

Correlation and covariance are symmetric, regression isn't. (Switching x and y axes changes the slope and intercept.)

$$s_{mpg,hp} = -321 \quad s_{mpg} = 6 \quad s_{hp} = 69 \quad \bar{mpg} = 20, \quad \bar{hp} = 147$$

What is the sample correlation between MPG (x) and horsepower (y)?

Horsepower and MPG

# Checking Our Results

$$s_{mpg,hp} = -321 \quad s_{mpg} = 6 \quad s_{hp} = 69 \quad \bar{mpg} = 20, \quad \bar{hp} = 147$$

What is the sample correlation
between MPG (x) and horsepower
(y)?

Horsepower and MPG

$$r = \frac{s_{xy}}{s_x s_y} = \frac{-321}{6 \times 69} \approx -0.78$$

$$s_{mpg,hp} = -321 \quad s_{mpg} = 6 \quad s_{hp} = 69 \quad \bar{mpg} = 20, \quad \bar{hp} = 147$$

What is the value of *b* for the regression:

$$\hat{y} = a + bx$$

where *x* is MPG and *y* is horsepower?

Horsepower and MPG

$$s_{mpg,hp} = -321 \quad s_{mpg} = 6 \quad s_{hp} = 69 \quad \bar{mpg} = 20, \quad \bar{hp} = 147$$

What is the value of b for the
regression:

$$\hat{y} = a + bx$$

where x is MPG and y is horsepower?

Horsepower and MPG

$$b = \frac{s_{xy}}{s_x^2} = \frac{-321}{6^2} \approx -8.9$$

$s_{mpg,hp} = -321$ $s_{mpg} = 6$ $s_{hp} = 69$ $\bar{mpg} = 20$, $\bar{hp} = 147$

What is the value of a for the regression:

$$\hat{y} = a + bx$$

where x is MPG and y is horsepower?
(prev. slide $b = -8.9$)

Horsepower and MPG

$s_{mpg,hp} = -321$ $s_{mpg} = 6$ $s_{hp} = 69$ $\bar{mpg} = 20,$ $\bar{hp} = 147$

What is the value of a for the
regression:

$$\hat{y} = a + bx$$

where x is MPG and y is horsepower?
(prev. slide $b = -8.9$)

Horsepower and MPG

$$a = \bar{y} - b\bar{x} = 147 - (-8.9) \times 20 \approx 325$$

Horsepower and MPG

$$a \approx 325$$

$$b \approx -8.9$$

$$\widehat{hp} = 324.08 - 8.83mpg$$

- Regression, Covariance, and Correlation: linear association.
- Linear association $\neq$ causation.
- Linear is not the only kind of association!

Correlation = 0

Please do the following before our next class:

1. Complete the "Odd Questions" quiz posted on Piazza
   `OddQuestions.pdf` – we'll be discussing these in class.

2. If you're rusty on permutations, combinations, etc. from
   High School math, read this review `http://ditraglia.com/`
   `Econ103Public/ClassicalProbability.pdf`

Review

Essential Distinction: Parameter vs. Statistic

N individuals in the Population, n individuals in the Sample:

	Parameter (Population)	**Statistic** (Sample)
Mean	$\mu_x = \dfrac{1}{N}\displaystyle\sum_{i=1}^{N} x_i$	$\bar{x} = \dfrac{1}{n}\displaystyle\sum_{i=1}^{n} x_i$
Var.	$\sigma_x^2 = \dfrac{1}{N}\displaystyle\sum_{i=1}^{N}(x_i - \mu)^2$	$s_x^2 = \dfrac{1}{n-1}\displaystyle\sum_{i=1}^{n}(x_i - \bar{x})^2$
S.D.	$\sigma_x = \sqrt{\sigma_x^2}$	$s_x = \sqrt{s^2}$
Cov.	$\sigma_{xy} = \dfrac{\sum_{i=1}^{N}(x_i - \mu_x)(y_i - \mu_y)}{N}$	$s_{xy} = \dfrac{\sum_{i=1}^{n}(x_i - \bar{x})(y_i - \bar{y})}{n-1}$
Corr.	$\rho = \dfrac{\sigma_{xy}}{\sigma_x \sigma_y}$	$r = \dfrac{s_{xy}}{s_x s_y}$

- Wonnacott: Chapter 2, Section 4-5 A+B, Section 5-3, 11-1, 11-2, Appendices to 2-2, 2-5, 11-1, and 11-2
- How to Lie with Statistics: Chapters 2, 5, and 6